

Lab course Image Analysis I ST 2023

Hubert Kanyamahanga
(kanyamahanga@ipi.uni-hannover.de)

Lab 1: Bayesian Classifiers



Content of the 1st Lab

Goal:

- Automatic extraction of **features** from images for **classification** task
 - By **probabilistic** image interpretation → **requires models/classifiers**

Tasks:

1. Implement your **methods to extract features** from images
2. Implement and train Bayesian classifiers to
 - **classify** real images using **extracted features**
3. Generate **synthetic datasets drawn** from Gaussian distribution and
 - **apply** the implemented classifiers on that toy dataset
4. Discussion and evaluation (**Visually** and **Quantitatively**)

Overview: Lab 1

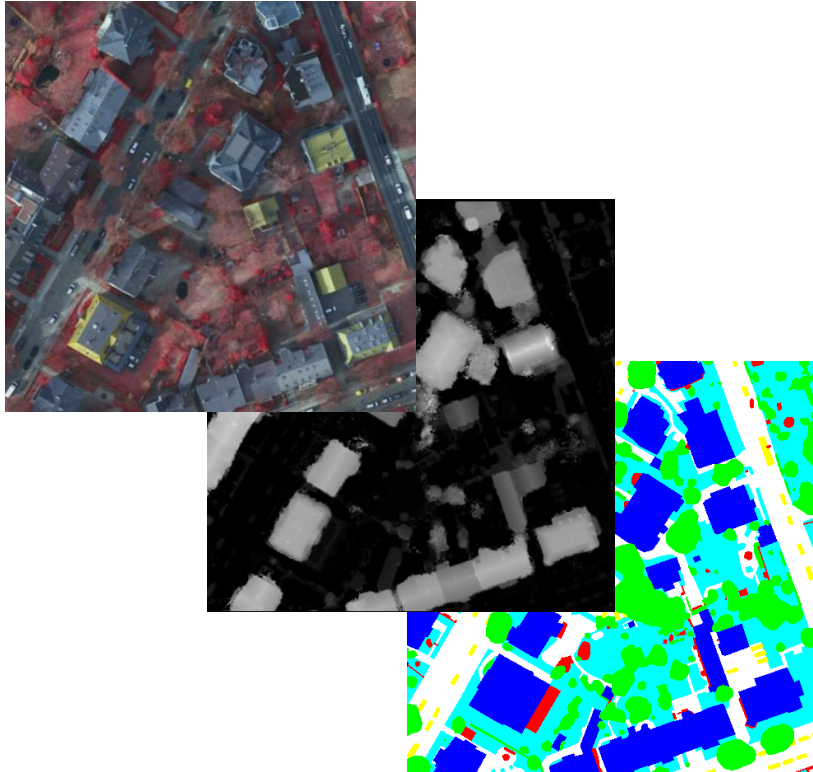
- Hand-crafted Features
- Bayesian Classification
- Generative Probabilistic Classifiers:
 - Normal Distribution (Gaussian) Classifier
 - Gaussian Mixture Model Classifier
- Quality Metrics

Feature Extraction

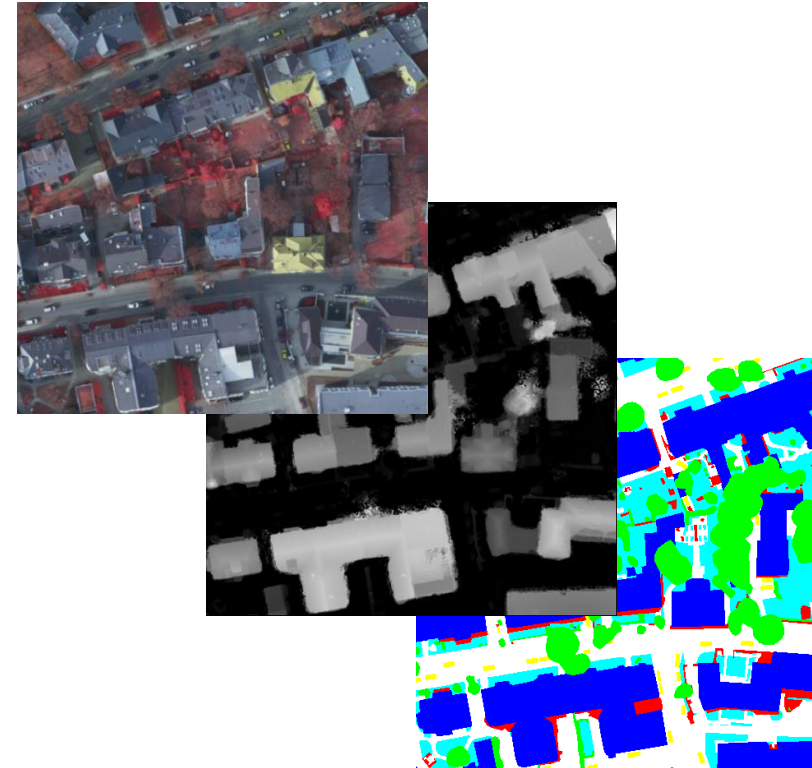
- Two patches (training / testing) with:
 - Channels: **Near infrared**, **red** and **green** (NIR,R,G)
 - Normalized **Digital Surface Model** (NDSM)
 - **Reference labels** (per pixel, 5 classes)
 - Classes are: 'STREET', 'HOUSE', 'LOW VEG.', 'HIGH VEG.', 'CAR'
- **Additional feature**: $NDVI = \frac{NIR-R}{NIR+R}$ (if $NIR = R = 0 \rightarrow NDVI = 0$) Note that $-1 < NDVI < +1$
- We will not use class, **clutter** → merge with class, **low vegetation**
- Normalization of channels required!
 - Shift and scale IR, R and G so that 0 will be mapped to -1.0 and 255 will be mapped to 1.0

Dataset

Training



Testing



Objective: Classification Task: USE extracted features, fit a probabilistic classifier

Overview: Lab 1

- Hand-crafted Features
- **Bayesian Classification**
- Generative Probabilistic Classifiers:
 - Normal Distribution (Gaussian) Classifier
 - Gaussian Mixture Model Classifier
- Quality Metrics



Bayesian Classification

Classification task: What is the class label L for a feature \mathbf{x} ?

MAP-Criterion: Classification based on maximum of the **posterior probability** $p(C | \mathbf{x})$

Generative approach:

- Often easier to model the **causal relation** between object type and observed features: the observed features are a function of the object type
- $p(C | \mathbf{x})$ is modelled indirectly according to the **Theorem of Bayes:**

$$p(C | \mathbf{x}) = \frac{p(\mathbf{x} | C) \cdot p(C)}{p(\mathbf{x})}$$

- The theorem of Bayes allows **inverse reasoning**: derive information about the **cause** (the class) from the **effect** (the observed features).

Bayesian Classification

- $p(\mathbf{x}|C)$: **Likelihood**

$$p(C | \mathbf{x}) = \frac{p(\mathbf{x} | C) \cdot p(C)}{p(\mathbf{x})}$$

- Probability to observe \mathbf{x} if it is known to belong to class C
- For each class C^k there is a model for $p(\mathbf{x}|C=C^k)$ that describes the distribution of the features for this class
- Determined from training data, but ?????? HOW ??????
- ~~Non-parametric Models: direct determination of $p(\mathbf{x} | C)$ from the training data~~
- **Parametric Models:** Based on the assumption of an analytical model for $p(\mathbf{x} | C)$ whose parameters are estimated from the training data → Machine Learning

Bayesian Classification

- $p(C)$: **Prior probability**

$$p(C | \mathbf{x}) = \frac{p(\mathbf{x} | C) \cdot p(C)}{p(\mathbf{x})}$$

- Corresponds to bias for the occurrence of C
- If no information is available: **Uniform Distribution** (e.g for 5 classes: $p(C1)=1/5$, $p(C2)=1/5$, $p(C3)=1/5$, $p(C4)=1/5$, $p(C5)=1/5$)
→ MAP becomes **Maximum Likelihood** (ML) criterion
- $p(C)$ can be determined iteratively :
 - 1) Classification under the assumption of a uniform distribution of the occurrence of the individual classes
 - 2) Determination of $p(C)$ from the relative frequencies of occurrence of the individual classes C^k
 - 3) Classification according to the theorem of Bayes



Bayesian Classification

- $p(\mathbf{x})$: **Probability of the data**

- Equal for all values of C because it does not depend on C

⇒ MAP can also be applied without knowing $p(\mathbf{x})$:

$p(C|\mathbf{x}) \propto p(\mathbf{x}|C) \cdot p(C)$ implies that

$$\max(p(C|\mathbf{x})) = \max(p(\mathbf{x}|C) \cdot p(C))$$

- $p(\mathbf{x})$ ensures that $p(C|\mathbf{x})$ can be **interpreted as a probability** and can be used as such in further probabilistic processes
- $p(\mathbf{x})$ can be determined as the **marginal distribution of $p(\mathbf{x}, C)$: $p(\mathbf{x})$ independent of C**

$$p(C|\mathbf{x}) = \frac{p(\mathbf{x}|C) \cdot p(C)}{p(\mathbf{x})}$$

$$p(\mathbf{x}) = \sum_k p(\mathbf{x} | C = L^k) \cdot p(C = L^k)$$

Summary: Bayesian Classification

$$p(C | \mathbf{x}) = \frac{p(\mathbf{x} | C) \cdot p(C)}{p(\mathbf{x})}$$

- Given:
 - Models for the likelihoods $p(\mathbf{x}|C=L^k)$ of all classes L^k
 - Prior probabilities $p(C=L^k)$ of all classes L^k
 - ??Question??: A feature vector \mathbf{x} to be classified
- Wanted: class C_{map} of \mathbf{x} according to the MAP criterion: *Maximize posterior probability $p(C | \mathbf{x})$*
- Procedure: Compute the posterior and take the class label to be the index of the maximum (ref. slide 9)
 - 1) For all L^k : calculate $p(\mathbf{x}, C=L^k) = p(\mathbf{x}|C=L^k) \cdot p(C=L^k)$
 - 2) Calculate $p(\mathbf{x}) = \sum_k p(\mathbf{x} | C = L^k) \cdot p(C = L^k)$
 - 3) For all L^k : calculate $p(C=L^k | \mathbf{x}) = p(\mathbf{x}, C=L^k) / p(\mathbf{x})$
 - 4) C_{map} is the label L^k for which $p(C=L^k | \mathbf{x})$ is a maximum

Overview: Lab 1

- Hand-crafted Features
- Bayesian Classification
- Generative probabilistic classifiers:
 - Normal Distribution (Gaussian) Classifier
 - Gaussian Mixture Model Classifier
- Quality Metrics



Generative Probabilistic Classifiers

How to model the likelihood? → **Learning** from data!

$$p(C | \mathbf{x}) = \frac{p(\mathbf{x} | C) \cdot p(C)}{p(\mathbf{x})}$$

Parametric Methods:

- **Analytical model** for the probability density $p(\mathbf{x}|C)$ is assumed
- The probability density function $p(\mathbf{x}|C)$ also depends on **parameters θ** , i.e. $p(\mathbf{x}|C) = p(\mathbf{x}|C, \theta)$
- Parameters are learned from training data
→ Training samples are required for each class $C=L^k$ to determine the parameters θ_k of $p(\mathbf{x}|C=L^k, \theta_k)$
- **BUT which parametric models do we USE?**

~~(Non-parametric methods are another option, but not covered in this lab)~~



1. Single Gaussian Model

- Frequent assumption: **Multivariate normal distribution**

$$p(\mathbf{x} | C = L^k) = \frac{1}{(2\pi)^{D/2} \cdot \|\Sigma_k\|^{1/2}} \cdot e^{-\frac{1}{2}(\mathbf{x} - \mu_k)^T \cdot \Sigma_k^{-1} \cdot (\mathbf{x} - \mu_k)}$$

- Extension of Univariate case, \mathbf{x} here is a feature Vector
- Maximum (Log-) Likelihood estimation of the parameters:

$$\sum_i \ln p(\mathbf{x}_{ik} | \theta_k) \rightarrow \max$$

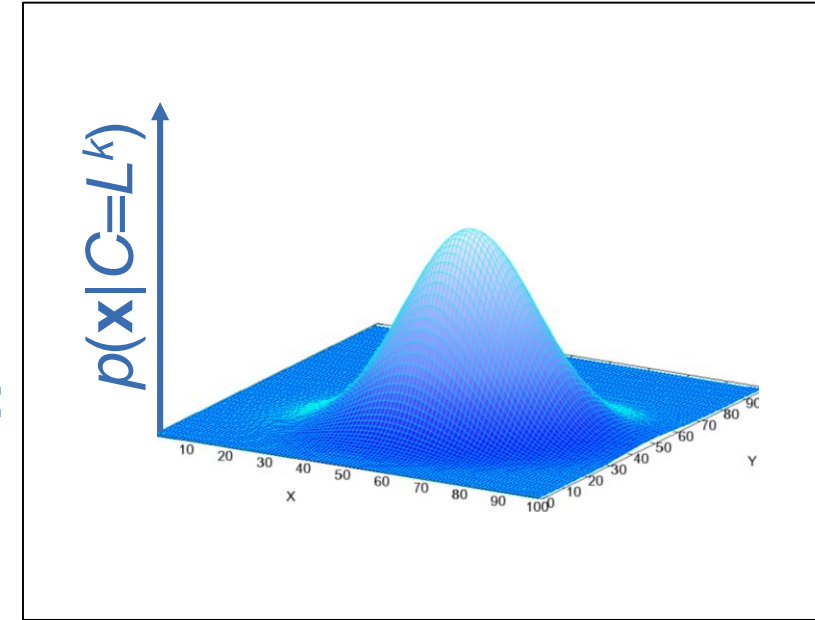
\Rightarrow Result for μ_k : $\mu_k = \frac{1}{N_k} \cdot \sum_i \mathbf{x}_{ik}$

\Rightarrow Result for Σ_k : $\Sigma_k = \frac{1}{N_k - 1} \cdot \sum_i (\mathbf{x}_{ik} - \mu_k) \cdot (\mathbf{x}_{ik} - \mu_k)^T$

$\theta_k == (\mu_k \text{ and } \Sigma_k)$:

Learned parameters for a
Single Gaussian Model

- Prerequisite: L^k must only correspond to one cluster in feature space (c./e.g. Slide 48/47)



Overview: Lab 1

- Hand-crafted Features
- Bayesian Classification
- Generative Probabilistic Classifiers:
 - Normal Distribution (Gaussian) Classifier
 - Gaussian Mixture Model Classifier
- Quality metrics



2. Gaussian Mixture Model

- In the case of N_j clusters for the class $C = L_k$, every cluster is described by a normal distribution

- The total probability density is obtained from the weighted sum of the components:

$$p(\mathbf{x} | C = L^k) = \sum_{j=1}^{N_j} \pi_j \cdot N(\mathbf{x} | \boldsymbol{\mu}_{kj}, \boldsymbol{\Sigma}_{kj})$$

with π_j ... Mixture coefficient for cluster j , corresponding to the prior probability for j

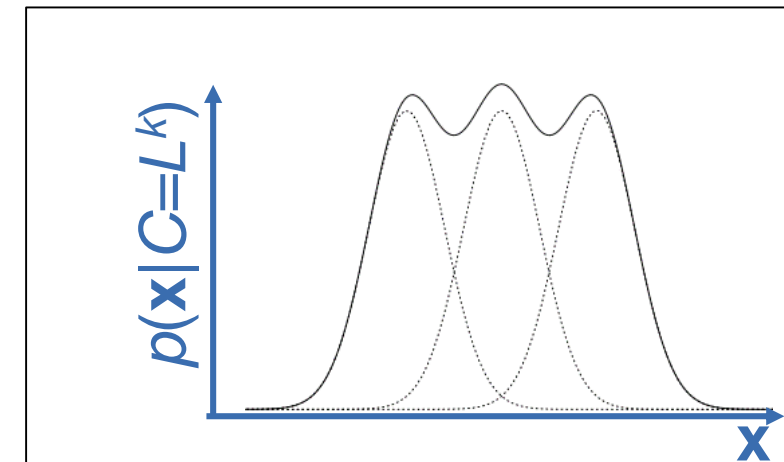
$\boldsymbol{\mu}_{kj}$... Mean value for cluster j

$\boldsymbol{\Sigma}_{kj}$... Covariance matrix for cluster j

$N(\mathbf{x} | \boldsymbol{\mu}_{kj}, \boldsymbol{\Sigma}_{kj})$... Probability density of the normal distribution for cluster j

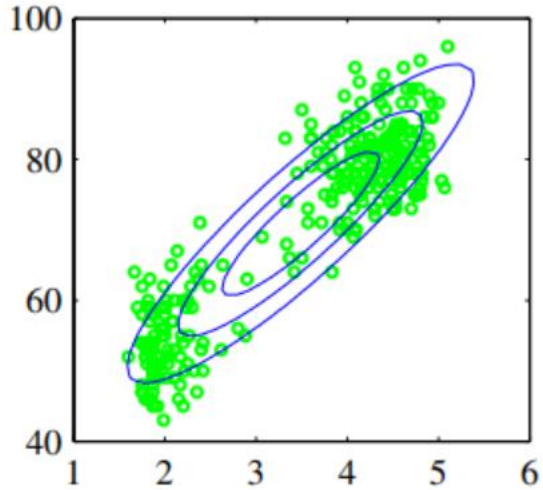
$\theta_k == (\pi_j, \mu_{kj} \text{ and } \Sigma_{kj})$:

Learned parameters for a Mixture of Gaussian Model

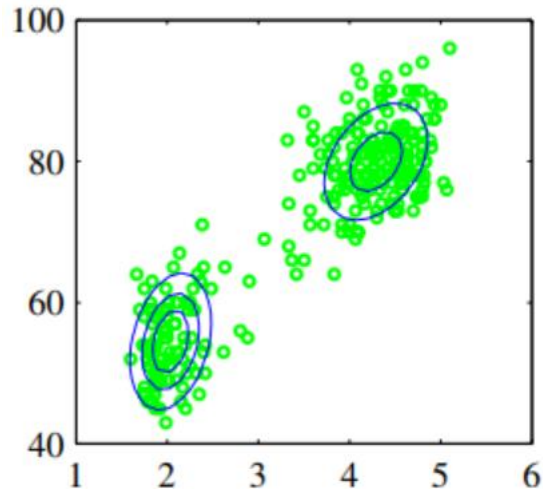


Gaussian Mixture Model: Example

Single Gaussian Model



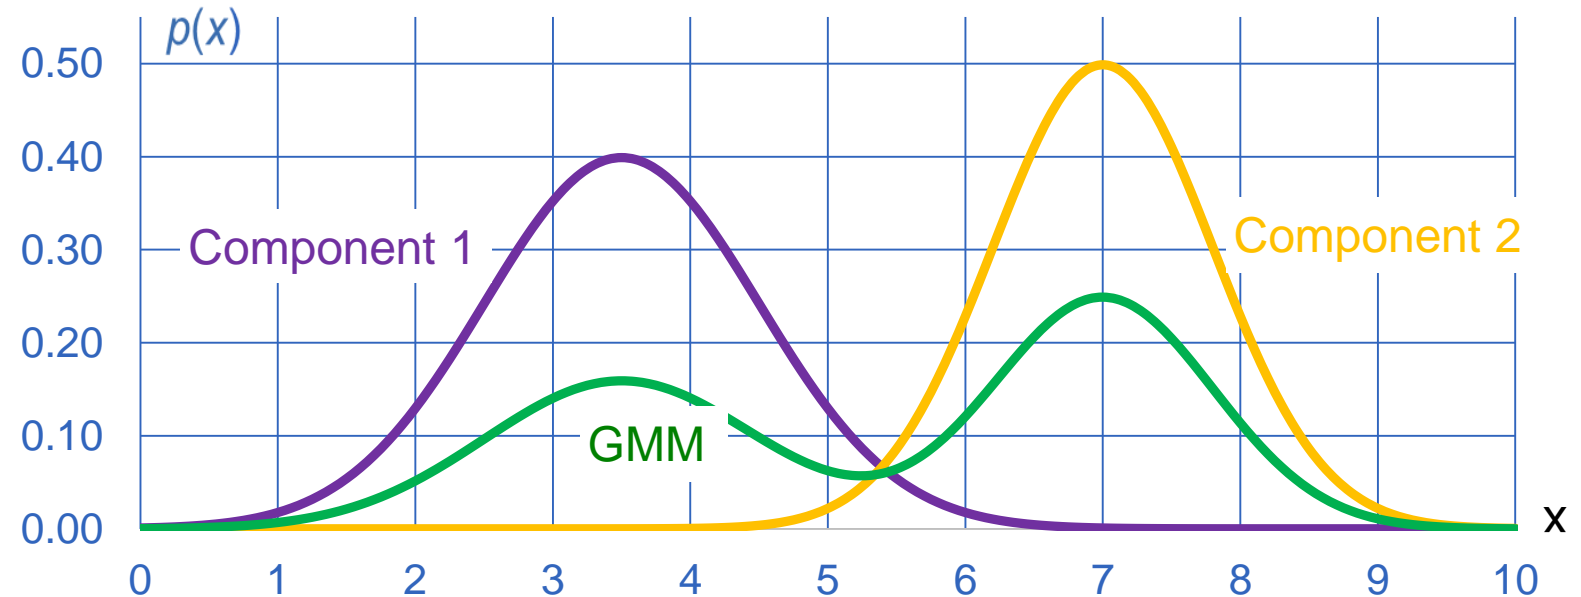
Gaussian Mixture Model



- One feature x , two Gaussian components (class index k is omitted):
 - Component 1 ($j = 1$): $\mu_1 = 3.5$, $\sigma_1 = 1.0$, $\pi_1 = 0.3$
 - Component 2 ($j = 2$): $\mu_2 = 7.0$, $\sigma_2 = 0.8$, $\pi_2 = 0.7$
 - GMM: $p(x) = \pi_1 \cdot N(x|\mu_1, \sigma_1) + \pi_2 \cdot N(x|\mu_2, \sigma_2)$

Linear combination of Gaussian distributions

- e.g.
Street
shadow/
sun



Summary: Gaussian Mixture Model

- **Issue:** We doesn't know which data points came from which latent component/cluster
- Parameters to be estimated: $\pi_j, \mu_{kj}, \Sigma_{kj}$ (on set of these params. per cluster j and class k)
- Training of the mixture model requires cluster analysis of the feature space → **unsupervised classification**
- Method: “**Expectation Maximisation**” (EM)
 - see lecture “Unsupervised Classification”
 - **E-step:** For each data point, compute the prob. of being generated by each component
 - **M-step:** Adjust the parameters to maximize the likelihood of the data given those params.
- EM requires the number of clusters N_j to be known in advance



Overview: Lab 1

- Hand-crafted Features
- Bayesian Classification
- Generative Probabilistic Classifiers:
 - Normal Distribution (Gaussian) Classifier
 - Gaussian Mixture Model Classifier
- Quality metrics



Quality Metrics

- **Quality metrics** are used to evaluate the performance of a classifier
- Evaluation requires a **reference** against which to compare the results of a classifier
- **Class-wise** metrics:
 - Precision (aka User's Accuracy, Correctness)
 - Recall (aka Producer's Accuracy, Completeness)
 - Intersection over Union (aka Quality)
 - F1 score
 - ...
- **Global** metrics:
 - Overall Accuracy
 - Mean F1 score
 - ...

[python script]

```
def calculate_quality_metrics(predicted, true_labels, classes):  
    TP = # True Positives  
    FP = # False Positives  
    FN = # False Negatives  
  
    p = TP/(TP+FP)  
    r = TP/(TP+FN)  
    f1 = 2*p*r/(p + r)  
  
    return p, r, f1
```

Quality Metrics

Terminology:

- True Positives (TP_k):
 - Number of features **correctly** classified as class k
- False Positives (FP_k):
 - Number of features **wrongly** classified as class k
- False Negatives (FN_k):
 - Number of features **wrongly** classified as **not** class k

Quality Metrics

- Precision:

$$P_k = \frac{TP_k}{TP_k + FP_k}$$

All samples classified as k

- Recall:

$$R_k = \frac{TP_k}{TP_k + FN_k}$$

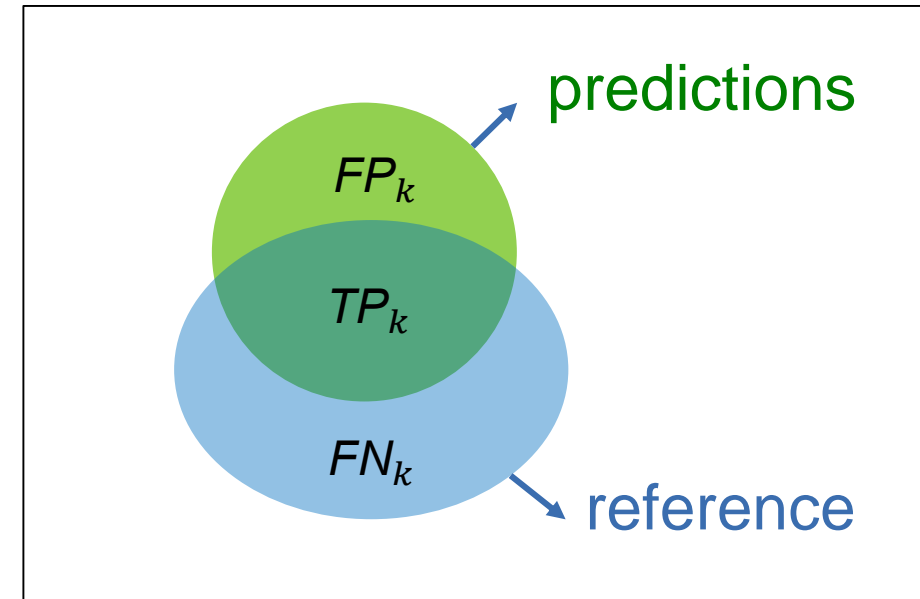
All samples that belong to class k

- I.o.U.:

$$IoU_k = \frac{TP_k}{TP_k + FP_k + FN_k}$$

- F1 score:

$$F_k = \frac{2 \cdot P_k \cdot R_k}{P_k + R_k}$$



Quality Metrics

- Overall accuracy:

$$OA = \frac{1}{N} \cdot \sum_{k=1}^M TP_k$$

- Mean F1 score:

$$mF1 = \frac{1}{M} \cdot \sum_{k=1}^M F_k$$

With:

- N : Number of samples
- M : Number of classes

Submission of Results

- Assignment

- Jupyter Notebook (only digital)

- Run every cell and save the notebook
- Use meaningful variable names (refer to [IAI_23_Lab_Python_Basics.ipynb](#))
- Write comments if required
- Answer concisely but completely
- Consider acceptance rules
- Consider [IAI_23_Lab_Technical_Details.pdf](#)



Submission of Results

- Submission deadline: **12. June 2023 before 11:00 am**
- There are **no resubmissions!**
- If you failed one or more labs in the last semester, you have to hand in **ALL** labs again!
- Introduction to the 2nd Lab: **12. June. 2023 at 1:00pm**

